

```

# Check GPU
!nvidia-smi
# install XGBoost for GPU support
!pip install xgboost

Thu Jan  1 16:26:03 2026
+-----+
| NVIDIA-SMI 550.54.15           Driver Version: 550.54.15
CUDA Version: 12.4      |
+-----+
| GPU Name             Persistence-M | Bus-Id     Disp.A |
Volatile Uncorr. ECC |
| Fan Temp Perf       Pwr:Usage/Cap |           Memory-Usage |
GPU-Util Compute M. |
|                         |                               |
MIG M. |
|
=====+=====+=====+=====
| 0 Tesla T4          Off | 00000000:00:04.0 Off |
0 |
| N/A    70C   P0        31W / 70W | 102MiB / 15360MiB |
0%      Default |
|                         |                               |
N/A |
+-----+
+-----+
| Processes:
|
| GPU GI CI      PID  Type  Process name
GPU Memory |
|           ID   ID
Usage      |
|
=====+=====
+-----+
Requirement already satisfied: xgboost in
/usr/local/lib/python3.12/dist-packages (3.1.2)
Requirement already satisfied: numpy in
/usr/local/lib/python3.12/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in
/usr/local/lib/python3.12/dist-packages (from xgboost) (2.27.5)

```

```
Requirement already satisfied: scipy in
/usr/local/lib/python3.12/dist-packages (from xgboost) (1.16.3)

import pandas as pd
import numpy as np
import time

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import accuracy_score

import xgboost as xgb
import matplotlib.pyplot as plt


df = pd.read_csv('/content/car_evaluation.csv', header=None)

df.columns = [
    "buying", "maint", "doors",
    "persons", "lug_boot", "safety", "target"
]

display(df.head())

{"summary": "\n  \"name\": \"display(df\", \n  \"rows\": 5,\n\"fields\": [\n    {\n      \"column\": \"buying\", \n      \"dtype\": \"category\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"vhigh\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    },\n    {\n      \"column\": \"maint\", \n      \"dtype\": \"category\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"vhigh\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    },\n    {\n      \"column\": \"doors\", \n      \"dtype\": \"category\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"2\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    },\n    {\n      \"column\": \"persons\", \n      \"dtype\": \"category\", \n      \"num_unique_values\": 1,\n      \"samples\": [\n        \"med\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    },\n    {\n      \"column\": \"lug_boot\", \n      \"dtype\": \"category\", \n      \"num_unique_values\": 2,\n      \"samples\": [\n        \"low\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    },\n    {\n      \"column\": \"safety\", \n      \"dtype\": \"string\", \n      \"num_unique_values\": 3,\n      \"samples\": [\n        \"low\"\n      ],\n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    }\n  },\n  {\n    \"column\": \"target\", \n    \"properties\": {\n      \"dtype\": \"category\", \n      \"num_unique_values\": 1\n    }\n  }\n]"}
```

```
1,\n      \"samples\": [\n          \"unacc\"\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    }\n  ]\n},\"type\":\"dataframe\"}\n\n# Count instances of each class\nlabel_counts = df["target"].value_counts()\n\nplt.figure(figsize=(6, 4))\nplt.bar(label_counts.index, label_counts.values)\nplt.title("Class Distribution of Car Evaluation Dataset")\nplt.xlabel("Car Acceptability Class")\nplt.ylabel("Number of Instances")\nplt.grid(axis='y')\nplt.show()
```



```
df.info()
df["target"].value_counts()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   buying      1728 non-null   object 
 1   maint       1728 non-null   object 
 2   doors        1728 non-null   object 
 3   horsepower  1728 non-null   int64  
 4   weight      1728 non-null   float64
 5   acceleration 1728 non-null   float64
 6   mpg         1728 non-null   float64
```

```
3    persons    1728 non-null    object
4    lug_boot   1728 non-null    object
5    safety     1728 non-null    object
6    target     1728 non-null    object
dtypes: object(7)
memory usage: 94.6+ KB

target
unacc    1210
acc      384
good     69
vgood    65
Name: count, dtype: int64

from sklearn.preprocessing import LabelEncoder

X = df.drop("target", axis=1)
y = df["target"]

# Encode target labels to numerical values
le = LabelEncoder()
y = le.fit_transform(y)

print("Encoded target labels:", le.classes_)

Encoded target labels: ['acc' 'good' 'unacc' 'vgood']

encoder = OneHotEncoder()

X_encoded = encoder.fit_transform(X).toarray()

print("Encoded feature shape:", X_encoded.shape)

Encoded feature shape: (1728, 21)

X_train, X_test, y_train, y_test = train_test_split(
    X_encoded,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

xgb_cpu = xgb.XGBClassifier(
    objective="multi:softprob",
    eval_metric="mlogloss",
    n_estimators=400,
    max_depth=6,
    learning_rate=0.1,
    subsample=0.9,
    colsample_bytree=0.9,
```

```

        tree_method="hist",    # CPU
        random_state=42
    )

start_cpu = time.time()
xgb_cpu.fit(X_train, y_train)
cpu_time = time.time() - start_cpu

cpu_preds = xgb_cpu.predict(X_test)
cpu_acc = accuracy_score(y_test, cpu_preds)

cpu_time, cpu_acc
(0.8878271579742432, 0.9971098265895953)

xgb_gpu = xgb.XGBClassifier(
    objective="multi:softprob",
    eval_metric="mlogloss",
    n_estimators=400,
    max_depth=6,
    learning_rate=0.1,
    subsample=0.9,
    colsample_bytree=0.9,
    tree_method="auto",   # Let XGBoost automatically detect GPU
    random_state=42
)

start_gpu = time.time()
xgb_gpu.fit(X_train, y_train)
gpu_time = time.time() - start_gpu

gpu_preds = xgb_gpu.predict(X_test)
gpu_acc = accuracy_score(y_test, gpu_preds)

gpu_time, gpu_acc
(0.5092513561248779, 0.9971098265895953)

print("Model      | Device | Time (s) | Accuracy")
print("-----")
print(f"XGBoost    | CPU    | {cpu_time:.4f} | {cpu_acc:.4f}")
print(f"XGBoost    | GPU    | {gpu_time:.4f} | {gpu_acc:.4f}")
print(f"\nSpeedup (CPU/GPU): {cpu_time / gpu_time:.2f}x")


Model      | Device | Time (s) | Accuracy
-----
XGBoost    | CPU    | 0.8878  | 0.9971
XGBoost    | GPU    | 0.5093  | 0.9971

Speedup (CPU/GPU): 1.74x

```

```
accuracies = [cpu_acc, gpu_acc]

plt.bar(models, accuracies)
plt.title("Accuracy Comparison")
plt.ylabel("Accuracy")
plt.ylim(0.9, 1.0)
plt.show()
```



```
speedup = cpu_time / gpu_time

plt.figure(figsize=(12, 5))

#Training Time Comparison
plt.subplot(1, 2, 1)
devices = ["CPU", "GPU"]
times = [cpu_time, gpu_time]

plt.plot(devices, times, marker='o', linewidth=3)
plt.fill_between(devices, times, alpha=0.2)
plt.title("Training Time: CPU → GPU")
plt.xlabel("Device")
plt.ylabel("Time (seconds)")
```

```


plt.grid(True)

#Speed-Up Ratio
plt.subplot(1, 2, 2)

plt.axhline(y=1, linestyle='--', linewidth=2)
plt.plot(["XGBoost"], [speedup], marker='o', markersize=10)
plt.title("CPU to GPU Speed-Up")
plt.ylabel("Speed-Up (CPU / GPU)")
plt.grid(True)

plt.tight_layout()
plt.show()

```



```

### Conclusion
# XGBoost achieved near-perfect accuracy on the Car Evaluation
dataset.
# GPU training was significantly faster than CPU training.
# Accuracy remained almost identical on both CPU and GPU.
# This demonstrates the benefit of parallel computing using GPU
acceleration.

```